

Final Lecture: Infinite Relational Model

Jesse Mu

2016-11-02

Contents

1	The Infinite Relational Model	1
1.1	Demo	1
1.1.1	Sanity check	1

1 The Infinite Relational Model

What happens when you want to cluster your data, but the number of clusters is unknown? While some approaches involve fitting several models with a varying number of clusters to your data and comparing model fit statistics, the Bayesian approach is to specify a model which is allowed to dynamically grow the number of clusters as the complexity of the data warrants.

The **Infinite Relational Model** is the prototypical example of such a model (Kemp et al. 2006).

For this last project, a simple version of Charles Kemp's Infinite Relational Model (IRM) was coded in `irm.R` to co-cluster rows and columns of a simple 2-dimensional binary relation.


1.1 Demo

```
source('irm.R')
```

1.1.1 Sanity check

As a sanity check, the toy matrix from the original paper is used:

```
R = rbind(
  c(0, 0, 1, 0, 1, 0, 0, 1, 0),
  c(0, 0, 0, 0, 0, 0, 1, 0, 1),
  c(0, 0, 1, 0, 0, 0, 1, 0, 1),
  c(0, 1, 1, 0, 0, 0, 0, 1, 1),
  c(0, 0, 0, 0, 0, 0, 1, 0, 1),
  c(0, 1, 1, 0, 1, 0, 0, 1, 0),
  c(1, 0, 0, 0, 0, 1, 0, 0, 0),
  c(0, 0, 0, 0, 0, 1, 1, 0, 1),
  c(1, 0, 0, 1, 0, 1, 0, 0, 0)
)
plot.R(R)
```



```

Z = irm(R, sweeps = 1000)
top.n(Z)
#> 122121323 122321424 122123424 122121324 122323424 123121424 122323454 122324525
#>      878      51      31      14      14      8      2      1
#> 123234565
#>      1
plot.R(R, mode.irm(Z))

```



This successfully finds the clusters of rows and columns that correspond to the original paper.

Kemp, Charles, Joshua B. Tenenbaum, Thomas L. Griffiths, Takeshi Yamada, and Naonori Ueda. 2006. "Learning Systems of Concepts with an Infinite Relational Model." *Proceedings of the National Conference on Artificial Intelligence (AAAI)* 21: 381–88. <http://web.mit.edu/cocosci/archive/Papers/Kemp-et al-AAAI06.pdf>.