

Final Lecture: Infinite Relational Model

Jesse Mu

2016-11-02

Contents

1	The Infinite Relational Model	1
1.1	Demo	1
1.1.1	Sanity check	1

1 The Infinite Relational Model

What happens when you want to cluster your data, but the number of clusters is unknown? While some approaches involve fitting several models with a varying number of clusters to your data and comparing model fit statistics, the Bayesian approach is to specify a model which is allowed to dynamically grow the number of clusters as the complexity of the data warrants.

The **Infinite Relational Model** is the prototypical example of such a model.

Kemp et al. (2006). Learning Systems of Concepts with an Infinite Relational Model¹

For this last project, I coded up a simple version of Charles Kemp's Infinite Relational Model (IRM) in `irm.R` to co-cluster rows and columns of a simple 2-dimensional binary relation.

1.1 Demo


```
source('irm.R')
```

1.1.1 Sanity check

As a sanity check, we use the toy matrix in the original paper:

```
R = rbind(
  c(0, 0, 1, 0, 1, 0, 0, 1, 0),
  c(0, 0, 0, 0, 0, 0, 1, 0, 1),
  c(0, 0, 1, 0, 0, 0, 1, 0, 1),
  c(0, 1, 1, 0, 0, 0, 0, 1, 1),
  c(0, 0, 0, 0, 0, 0, 1, 0, 1),
  c(0, 1, 1, 0, 1, 0, 0, 1, 0),
  c(1, 0, 0, 0, 0, 1, 0, 0, 0),
  c(0, 0, 0, 0, 0, 1, 1, 0, 1),
  c(1, 0, 0, 1, 0, 1, 0, 0, 0)
)
plot.R(R)
```

¹<http://web.mit.edu/cocosci/Papers/Kemp-et-al-AAAI06.pdf>



```
Z = irm(R, sweeps = 1000)
top.n(Z)
#> 122121323 122121324 122321424 122123424 123121424 122121343 122123454 122324525
#>      903      35      19      13      13      6      3      3
#> 123124535 122321456
#>      2      1
plot.R(R, mode.irm(Z))
```



This successfully finds the clusters of rows and columns that correspond to the original paper.